

Legal Clause Similarity Analysis

Deep Learning Assignment 2 - CS452

Student Name: Abdul Rehman

Roll Number: 21I-1780

Date: November 10, 2025

Table of Contents:

Legal Clause Similarity Analysis.....	1
Deep Learning Assignment 2 - CS452.....	1
Table of Contents:.....	2
Executive Summary.....	4
1. Introduction.....	4
Problem Statement.....	4
2. Dataset Analysis.....	4
2.1 Dataset Overview.....	4
2.2 Dataset Splits.....	4
2.3 Class Distribution.....	5
2.4 Pair Generation Strategy.....	5
3. Preprocessing and Text Representation.....	5
3.1 Text Preprocessing Pipeline.....	5
3.2 Input Representation.....	5
4. Model Architectures.....	5
4.1 Baseline 1: BiLSTM Siamese Network.....	5
Architecture Details.....	5
Training Configuration.....	6
4.2 Baseline 2: Attention-Based Encoder.....	6
Architecture Details.....	6
Training Configuration.....	7
5. Training Results.....	7
5.1 BiLSTM Model Training.....	7
5.2 Attention Model Training.....	8
6. Performance Evaluation.....	9
6.1 Evaluation Metrics.....	9
6.2 BiLSTM Performance.....	9
6.3 Attention Model Performance.....	11
7. Comparative Analysis.....	12
7.1 Quantitative Comparison.....	12
7.2 Model Strengths and Weaknesses.....	12
BiLSTM Siamese Network.....	12
Attention-Based Encoder.....	13
7.3 Why BiLSTM Outperformed Attention.....	13
8. Qualitative Analysis.....	13
8.1 Correctly Classified Examples.....	13
Example 1: True Similar - Predicted Similar.....	13
Example 2: True Different - Predicted Different.....	14
Example 3: True Different - Predicted Different.....	14

8.2 Incorrectly Classified Example.....	14
Example 4: True Different - Predicted Similar.....	14
9. Domain-Specific Metric Analysis.....	15
9.1 Metric Selection for Legal Applications.....	15
Scenario 1: Contract Review & Deduplication.....	15
Scenario 2: Automatic Clause Merging.....	15
Scenario 3: Legal Research & Case Retrieval.....	15
9.2 Recommended Production Metric.....	15
10. Challenges and Limitations.....	16
10.1 Dataset Limitations.....	16
10.2 Model Limitations.....	16
10.3 Attention Model Challenges.....	16
11. Future Improvements.....	16
11.1 Model Enhancements.....	16
11.2 Architecture Variations.....	16
11.3 Training Strategies.....	17
12. Conclusion.....	17
Key Findings:.....	17
Practical Impact:.....	17
References.....	17
Appendix: Code Repository.....	18

Executive Summary

This report presents the implementation and evaluation of two baseline deep learning architectures for identifying semantic similarity between legal clauses. The task involves determining when two legal clauses convey the same or closely related meanings, despite differences in wording. The BiLSTM Siamese Network achieved outstanding performance with 94.59% accuracy, while the Attention-based Encoder demonstrated the challenges of training attention mechanisms from scratch.

1. Introduction

Legal documents utilize highly formal and structured language, with the same legal principle often expressed in multiple ways across different jurisdictions and contracts. This project addresses the challenge of automatically identifying semantic equivalence between legal clauses through deep learning, without relying on pre-trained transformer models.

Problem Statement

Given two legal clauses, determine whether they are semantically similar (express the same legal principle) or different (express distinct legal concepts).

2. Dataset Analysis

2.1 Dataset Overview

- **Source:** Legal Clause Dataset from Kaggle
- **Structure:** 395 category files containing legal clauses
- **Total Clauses:** 155,000+ individual legal clauses
- **Categories:** Various legal clause types including acceleration, arbitration, confidentiality, indemnification, governing law, etc.

2.2 Dataset Splits

After generating clause pairs for similarity classification:

Split	Samples	Percentage
Training	6,818	70%
Validation	1,461	15%
Testing	1,461	15%

2.3 Class Distribution

- **Positive Pairs (Similar):** 4,740 (48.7%)
- **Negative Pairs (Different):** 5,000 (51.3%)
- **Total Pairs:** 9,740

The dataset is well-balanced, making accuracy a reliable metric for evaluation.

2.4 Pair Generation Strategy

- **Positive Pairs:** Generated from clauses within the same category (assumed to be semantically similar)
- **Negative Pairs:** Generated from clauses across different categories (assumed to be semantically different)

3. Preprocessing and Text Representation

3.1 Text Preprocessing Pipeline

1. **Tokenization:** Convert text to sequences of integers using Keras Tokenizer
2. **Vocabulary Size:** 10,000 most frequent words
3. **Sequence Length:** Fixed at 200 tokens with post-padding/truncating
4. **Out-of-Vocabulary:** Handled with **<OOV>** token

3.2 Input Representation

- Each clause pair is represented as two parallel sequences
- Vocabulary: 19,431 unique tokens identified in the corpus
- Effective vocabulary used: 10,000 (for computational efficiency)

4. Model Architectures

4.1 Baseline 1: BiLSTM Siamese Network

Architecture Details

Rationale: Bidirectional LSTMs are well-suited for capturing sequential dependencies and contextual information in legal text. The Siamese architecture allows the model to learn a shared representation space where similar clauses are embedded closely.

Architecture Components:

1. Embedding Layer

- Input dimension: 10,000
- Output dimension: 128
- Mask zero: True (handles padding)

2. Encoder Network (Shared)

- Bidirectional LSTM Layer 1: 64 units (return sequences)
- Bidirectional LSTM Layer 2: 32 units
- Dense layer: 128 units with ReLU activation
- Dropout: 0.3

3. Similarity Computation

- Subtract: Element-wise difference between encodings
- Multiply: Element-wise product between encodings
- Concatenate: [encoding1, encoding2, difference, product]

4. Classification Head

- Dense: 64 units with ReLU
- Dropout: 0.3
- Dense: 32 units with ReLU
- Output: 1 unit with sigmoid activation

Total Parameters: 1,463,297 (5.58 MB)

Training Configuration

- **Optimizer:** Adam
- **Loss Function:** Binary Cross-Entropy
- **Batch Size:** 32
- **Epochs:** 30 (with early stopping)
- **Learning Rate:** 0.001 (with ReduceLROnPlateau)
- **Early Stopping:** Patience of 5 epochs on validation loss
- **Best Model:** Restored from epoch with lowest validation loss

4.2 Baseline 2: Attention-Based Encoder

Architecture Details

Rationale: Self-attention mechanisms can capture long-range dependencies and semantic relationships without the sequential bottleneck of RNNs. Multi-head attention allows the model to attend to different semantic aspects simultaneously.

Architecture Components:

1. Embedding Layer

- Input dimension: 10,000
- Output dimension: 128
- Mask zero: True

2. Encoder Network (Shared)

- Multi-Head Self-Attention: 4 heads, key dimension 32
- Add & LayerNormalization (residual connection)
- Feed-Forward Network: Dense(128, ReLU) → Dense(128)
- Add & LayerNormalization (residual connection)
- Global Average Pooling
- Dense: 128 units with ReLU
- Dropout: 0.3

3. Similarity Computation

- Same as BiLSTM: [encoding1, encoding2, difference, product]

4. Classification Head

- Dense: 64 units with ReLU
- Dropout: 0.3
- Dense: 32 units with ReLU
- Output: 1 unit with sigmoid

Total Parameters: 1,431,041 (5.46 MB)

Training Configuration

- Same hyperparameters as BiLSTM model
- Training stopped early due to validation performance plateau

5. Training Results

5.1 BiLSTM Model Training

The BiLSTM model demonstrated excellent learning characteristics:


Training Progression:

- **Epoch 1:** 50.35% accuracy → validation rapidly improved to 53.73%
- **Epoch 2:** Major breakthrough - validation accuracy jumped to 89.80%
- **Epoch 3-7:** Steady improvement, reaching 94.05% validation accuracy
- **Epoch 8-12:** Fine-tuning with learning rate reduction
- **Best Epoch:** 7 with validation accuracy of 94.05%

Convergence:

- Training loss decreased from 0.6935 to 0.0669
- Validation loss reached minimum of 0.1764 at epoch 7
- No significant overfitting observed (validation performance tracked training closely)

Training Graphs for BiLSTM:



5.2 Attention Model Training

The Attention model showed slower convergence:


Training Progression:

- **Epoch 1-2:** Struggling to learn (50-55% accuracy)
- **Epoch 3-4:** Gradual improvement (56-66% accuracy)
- **Epoch 5-9:** Continued learning but validation plateaued at ~64-65%
- Learning rate reduced at epoch 8, but minimal improvement

Challenges Observed:

- Much slower convergence compared to BiLSTM
- Validation loss increased while training loss decreased (overfitting signs)
- Stopped at epoch 9 due to early stopping criteria

Training Graphs for Attention Model:



6. Performance Evaluation

6.1 Evaluation Metrics

The following metrics were used to evaluate model performance:


Metric	Description	Importance for Legal Clause Similarity
Accuracy	Overall correct classifications	Useful given balanced dataset
Precision	True similar / Predicted similar	Critical - false positives mean incorrectly merging different clauses
Recall	True similar / Actually similar	Critical - false negatives mean missing equivalent clauses
F1-Score	Harmonic mean of Precision & Recall	Balanced measure for legal applications
ROC-AUC	Ranking ability across thresholds	Measures discriminative power
PR-AUC	Precision-Recall curve area	Better for slightly imbalanced scenarios

6.2 BiLSTM Performance

Metric	Score
Accuracy	94.59%
Precision	90.20%
Recall	99.72%

F1-Score	94.72%
ROC-AUC	97.22%
PR-AUC	94.26%

Confusion Matrix:



	Predicted Different	Predicted Similar
Actually Different	672 (90%)	78 (10%)
Actually Similar	2 (0.3%)	709 (99.7%)


Key Observations:

- Exceptional recall (99.72%) - rarely misses similar clauses
- Good precision (90.20%) - some false positives but acceptable
- High ROC-AUC (97.22%) - excellent ranking capability
- The model errs on the side of caution, preferring false positives over false negatives

6.3 Attention Model Performance

Metric	Score
Accuracy	65.71%
Precision	64.54%
Recall	65.54%
F1-Score	65.04%
ROC-AUC	71.77%
PR-AUC	66.13%

Confusion Matrix:




	Predicted Different	Predicted Similar
Actually Different	494 (66%)	256 (34%)
Actually Similar	245 (34%)	466 (66%)

Key Observations:

- Balanced but poor performance across all metrics
- Similar precision and recall (both ~65%)
- Low ROC-AUC (71.77%) indicates weak discriminative ability
- Model struggles to learn meaningful representations

7. Comparative Analysis

7.1 Quantitative Comparison



Model	Accuracy	Precision	Recall	F1-Score	ROC-AUC	Training Time
BiLSTM	94.59%	90.20%	99.72%	94.72%	97.22%	~12s/epoch
Attention	65.71%	64.54%	65.54%	65.04%	71.77%	~3s/epoch

Performance Gap: The BiLSTM outperforms the Attention model by approximately 29 percentage points in accuracy.

7.2 Model Strengths and Weaknesses

BiLSTM Siamese Network

Strengths:

1. **Superior Performance:** Achieved >94% across all major metrics
2. **Excellent Recall:** 99.72% recall means it catches nearly all similar clauses
3. **Stable Training:** Consistent convergence with no significant overfitting
4. **Robust Representations:** BiLSTM effectively captures sequential legal text patterns
5. **Practical Value:** High reliability for real-world legal document analysis

Weaknesses:

1. **Computational Cost:** Slower training (~12s/epoch vs 3s/epoch)
2. **Sequential Processing:** Cannot parallelize sequence processing
3. **False Positives:** 10% of different clauses incorrectly marked as similar
4. **Memory Requirements:** Larger model size (5.58 MB)

Attention-Based Encoder

Strengths:

1. **Faster Training:** 4x faster per epoch than BiLSTM
2. **Parallelizable:** Attention mechanism allows parallel computation
3. **Theoretical Advantages:** Can capture long-range dependencies more directly
4. **Smaller Model:** Slightly fewer parameters (5.46 MB)

Weaknesses:

1. **Poor Performance:** Only 65.71% accuracy - barely better than random
2. **Training Instability:** Showed signs of overfitting early
3. **Insufficient Convergence:** Needed more epochs or different architecture
4. **High Error Rate:** 34% false positive and false negative rates
5. **Not Production-Ready:** Cannot be deployed in current state

7.3 Why BiLSTM Outperformed Attention

Several factors contributed to the performance gap:

1. **Data Efficiency:** BiLSTMs are more sample-efficient for smaller datasets
2. **Inductive Bias:** Sequential processing aligns well with natural language structure
3. **Training Stability:** Recurrent connections provide more stable gradients in this setting
4. **Architecture Maturity:** BiLSTM is a well-established architecture for text
5. **Attention Complexity:** Attention mechanisms may require more data/tuning to excel

8. Qualitative Analysis

8.1 Correctly Classified Examples

Example 1: True Similar - Predicted Similar

Clause 1: "Interpretation. The article and section headings contained in this Agreement are inserted for convenience..."

Clause 2: "Interpretation. The words 'include,' 'includes,' and 'including' shall be deemed to be followed by..."

Prediction: SIMILAR (Confidence: 99.3%)

Analysis: Both clauses deal with interpretation rules for the agreement. The model correctly identified semantic similarity despite different specific content.

Example 2: True Different - Predicted Different

Clause 1: "Renewal Option. Tenant shall have the option, upon six (6) months' prior written notice to Landlord..."

Clause 2: "Performance. SCC will use its best efforts to achieve the goals of the FUND but SCC is not guaranteed..."

Prediction: DIFFERENT (Confidence: 100%)

Analysis: Clear semantic difference - one about lease renewal, other about performance obligations. Model correctly distinguished.

Example 3: True Different - Predicted Different

Clause 1: "Indemnification by the Company. The Company agrees to indemnify and hold harmless the Sales Agent..."

Clause 2: "Notice of Defaults. If a Default occurs and is continuing and if it is either actually known..."

Prediction: DIFFERENT (Confidence: 100%)

Analysis: Distinct legal concepts (indemnification vs. default notice). Model showed strong discriminative ability.

8.2 Incorrectly Classified Example

Example 4: True Different - Predicted Similar

Clause 1: "Environmental Laws. The Company (i) is in compliance with any and all Environmental Laws applicable..."

Clause 2: "Condemnation. All proceeds of condemnation awards or proceeds of sale in lieu of condemnation..."

Prediction: SIMILAR (Confidence: 94.1%)

Analysis: The model incorrectly classified these different clauses as similar. This is a false positive - both clauses involve legal compliance/obligations, which may have triggered similar embeddings. This highlights the need for deeper semantic understanding beyond surface-level patterns.

9. Domain-Specific Metric Analysis

9.1 Metric Selection for Legal Applications

For a legal clause similarity system deployed in production, the choice of primary metric depends on the use case:

Scenario 1: Contract Review & Deduplication

Priority: High Recall (minimize false negatives)

- **Rationale:** Missing a similar clause could mean overlooking critical legal obligations
- **Cost of Error:** Legal liability, contract disputes
- **Recommended Threshold:** Lower threshold to catch all potential similarities
- **Best Model:** BiLSTM with 99.72% recall

Scenario 2: Automatic Clause Merging

Priority: High Precision (minimize false positives)

- **Rationale:** Incorrectly merging different clauses could create legal conflicts
- **Cost of Error:** Invalid contract terms, legal vulnerabilities
- **Recommended Threshold:** Higher threshold to ensure confident matches only
- **Trade-off:** BiLSTM's 90.20% precision may need threshold adjustment

Scenario 3: Legal Research & Case Retrieval

Priority: ROC-AUC (ranking quality)

- **Rationale:** Users need a ranked list of similar clauses to review
- **Cost of Error:** Wasted time reviewing irrelevant clauses
- **Recommended Metric:** ROC-AUC (97.22% indicates excellent ranking)
- **Best Model:** BiLSTM with superior discriminative power

9.2 Recommended Production Metric

F1-Score should be the primary metric for general legal clause similarity systems because:

1. **Balanced Consideration:** Weighs both precision and recall equally
2. **Practical Value:** Reflects real-world performance across use cases
3. **Legal Requirements:** Legal applications need both high accuracy and completeness
4. **Risk Management:** Balances the costs of false positives and false negatives

Secondary Metrics:

- **ROC-AUC:** For evaluating ranking quality in search applications

- **Precision @ K:** For top-K retrieval scenarios
- **Recall:** For compliance-critical applications

10. Challenges and Limitations

10.1 Dataset Limitations

1. **Category-Based Similarity:** Assumption that same category = similar may not always hold
2. **No Fine-Grained Labels:** Binary similarity may miss nuanced relationships
3. **Limited Cross-Domain:** Clauses from specific legal domains may not generalize

10.2 Model Limitations

1. **No Pre-training:** Models trained from scratch lack broad language understanding
2. **Fixed Vocabulary:** OOV tokens may lose important legal terminology
3. **Sequence Length:** 200 tokens may truncate longer complex clauses
4. **No Legal Knowledge:** Models don't understand legal concepts explicitly

10.3 Attention Model Challenges

1. **Insufficient Training:** May need more epochs or better initialization
2. **Architecture Tuning:** Hyperparameters may not be optimal
3. **Data Efficiency:** Attention may require larger datasets to excel
4. **Training Stability:** More sensitive to hyperparameter choices

11. Future Improvements

11.1 Model Enhancements

1. **Ensemble Methods:** Combine BiLSTM and Attention predictions
2. **Hierarchical Models:** Sentence-level then document-level encoding
3. **External Embeddings:** Use GloVe or Word2Vec for initialization
4. **Data Augmentation:** Paraphrase clauses to increase training data

11.2 Architecture Variations

1. **CNN-LSTM Hybrid:** Combine convolutional and recurrent layers
2. **Deeper Networks:** Add more LSTM/Attention layers
3. **Better Attention:** Try Transformer-style multi-layer attention
4. **Regularization:** Experiment with different dropout rates

11.3 Training Strategies

1. **Learning Rate Scheduling:** More aggressive schedules
2. **Curriculum Learning:** Start with easy pairs, progress to hard ones
3. **Hard Negative Mining:** Focus on challenging negative examples
4. **Cross-Validation:** K-fold validation for more robust evaluation

12. Conclusion

This project successfully developed and evaluated two baseline architectures for legal clause similarity detection. The BiLSTM Siamese Network achieved excellent performance (94.59% accuracy, 97.22% ROC-AUC), demonstrating that recurrent architectures remain highly effective for legal text understanding even without pre-trained models.

Key Findings:

1. **BiLSTM Superiority:** Significantly outperformed attention-based approach
2. **High Recall Critical:** 99.72% recall ensures comprehensive similarity detection
3. **Practical Viability:** BiLSTM model ready for deployment with appropriate threshold tuning
4. **Attention Potential:** With more training/tuning, attention may close the gap

Practical Impact:

The developed system can:

- Automate contract review and reduce legal costs
- Identify redundant or conflicting clauses
- Support legal research and case law retrieval
- Enable intelligent document comparison tools

This work demonstrates that relatively simple architectures, when properly designed and trained, can achieve strong performance on complex legal NLP tasks without requiring massive pre-trained models.

References

1. Legal Clause Dataset: <https://www.kaggle.com/datasets/bahushruth/legalclausedataset>
2. Keras Documentation: <https://keras.io/>
3. Siamese Networks for Text Similarity (Bromley et al., 1993)
4. Bidirectional LSTM for Sequence Classification (Graves & Schmidhuber, 2005)
5. Attention Mechanisms in NLP (Bahdanau et al., 2014)

Appendix: Code Repository

GitHub Repository: <https://github.com/AbdulRehman1780/Legal-Clause-Similarity-Analysis.git>

The complete implementation includes:

- Modular data loading and preprocessing pipeline
- Two baseline model architectures
- Comprehensive evaluation framework
- Visualization utilities
- Training and inference scripts

Total Lines of Code: ~600 **Documentation:** Fully commented with docstrings **Code Quality:** Follows PEP 8 standards and OOP principles